

Canny Edge Detection

- Most widely used edge detector in computer vision.
- First derivative of the Gaussian closely approximates the operator that optimizes the product of *signal-to-noise* ratio and localization.
- Analysis based on "step-edges" corrupted by "additive Gaussian noise".

Least squares with binomial weights ~~~ edge detector.

J. Canny, *A Computational Approach To Edge Detection*.

IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.


Edge Detection Criteria

- Criteria for **optimal edge detection** (Canny 86):
 - Good detection accuracy:
 - minimize the probability of false positives (detecting spurious edges caused by noise),
 - false negatives (missing real edges)
 - Good localization:
 - edges must be detected as close as possible to the true edges.
 - Single response constraint:
 - minimize the number of local maxima around the true edge (i.e. detector must return single point for each true edge point)

- Examples... valid mostly for straight edges...




True
edge



Poor robustness
to noise



Poor
localization



Too many
responses

Canny Edge Detection

Steps:

- 1+2. Gaussian smoothing together with derivative of Gaussian (~discrete)
 3. Find magnitude and orientation of gradient
 4. Extract edge points: Non-maximum suppression
 5. Linking and thresholding: Hysteresis
- **MATLAB:** `edge(I, 'canny')`

First Two Steps

- Smoothing

$$I' = g(x, y) * I$$

$$g(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

Can be done with two one-dimensional filters.

- Derivative

$$S = \nabla(g * I) = (\nabla g) * I =$$

$$= \begin{bmatrix} g_x \\ g_y \end{bmatrix} * I = \begin{bmatrix} g_x * I \\ g_y * I \end{bmatrix}$$

$$\nabla g = \begin{bmatrix} \frac{\partial g}{\partial x} \\ \frac{\partial g}{\partial y} \end{bmatrix} = \begin{bmatrix} g_x \\ g_y \end{bmatrix}$$


$$h = g$$

two dimensional Gaussian


$$h_x(x,y) = \frac{\partial h(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$h_y(x,y) = \frac{\partial h(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

Scale



Example:



$$S = \nabla(g * I) = (\nabla g) * I$$

$$S = \begin{bmatrix} S_x & S_y \end{bmatrix} = \text{gradient vector}$$



sigma 1 pixel



2 pixels

Increased smoothing:

- Eliminates noise edges.
- Makes edges smoother and thicker.
- Removes fine detail.

Third Step

- magnitude and direction of $S = \begin{bmatrix} S_x & S_y \end{bmatrix}$

$$\text{magnitude} = \sqrt{(S_x^2 + S_y^2)}$$

$$\text{direction} = \theta = \tan^{-1} \frac{S_y}{S_x}$$



image





gradient magnitude

Non-maximum suppression along the
direction of gradient

Fourth Step

At q , we have a maximum if the value is larger than those at both p and r . Interpolate to get these values.



Example: Non-Maximum Suppression



Original image



Gradient magnitude



courtesy of G. Loy

Non-maxima
suppressed



high threshold
strong edges only



low threshold
weak edges too


Fifth Step: Hysteresis Thresholding

- Hysteresis: no LOW maybe HIGH sure.
- Maintain two thresholds k_{high} and k_{low}
 - Use k_{high} to find strong edges to start edge chain.
 - Use k_{low} to find weak edges along the edge chain.
- Typical ratio of thresholds is roughly

$$k_{\text{high}} / k_{\text{low}} = 2 - 2.5$$

Closing edge gaps

- Check that maximum value of gradient value is sufficiently large and...
... use hysteresis.
 - use a high threshold to start edge curves and a low threshold to continue them.



Example

Original
image



Strong
edges
only



gap is gone



Strong +
connected
weak edges

Weak
edges
too



courtesy of G. Loy

Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$



Canny with $\sigma = 2$

- The choice of σ depends on desired behavior
 - large σ detects large scale edges
 - small σ detects fine features

Example of Canny edge detection



original image (Lena)


Compute Gradients (DoG)



X-Derivative of Gaussian



Y-Derivative of Gaussian



Gradient Magnitude
and orientation

Before non-max suppression...



...after non-max suppression



Before the hysteresis thresholding

- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



Final Canny Edges



A hidden advantage for the human observer.
She/he first see the original image and only
after the edges detected.

What happens if she/he cannot see the original
image first and therefore can rely on it?

This is how all the computer vision algorithms
has to work... all the time!



